

# Python in the Geosciences

Fall 2015 Seminar - University of Washington

Time: 1st Tuesday of the month, 4:00PM with Happy Hour to follow.

Location: 6th floor of the Physics Tower, Seminar Room C607

# Python in the Geosciences

Other Workshops / Seminars / Resources:

**Atmos-Python Workshop:**

[atmos-python@uw.edu](mailto:atmos-python@uw.edu) - (Andre or Jeremy)

**UW eScience Python Seminar: Fall 2015:**





<http://uwescience.github.io/python-seminar-2015/seminar/>

**AGU Fall Meeting 2015 Session:**

Python Solutions for the Earth Sciences: IN041:

# Data analysis tools for modern Python: *pandas, xray, and beyond*

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$


xray

Joe Hamman

October 6, 2015

Python in the Geosciences

Fall 2015 Seminar - University of Washington







Department of Civil  
and Environmental  
Engineering

# Audience Participation (familiarity with...)



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



xray

# **We should stop rewriting the same things...**

Working Assumption:

It is better to push repetitive tasks to well tested, optimized packages (e.g. pandas/xray/NumPy)



IPython





SymPy



pandas

$$y_{it} = \beta x_{it} + \mu_i + \epsilon_{it}$$



xray



machine learning in Python



scikit-image  
image processing in python



NetworkX

# The core of scientific computing in Python




- N-dimensional array objects
- High level functions
- Linear algebra functions
- random number generation
- Foreign language (C/Fortran) integration



- Depends on NumPy
- High level science / engineering modules
  - Fourier transforms
  - Integration
  - Interpolation
  - Optimization
  - Signal processing
  - Statistics

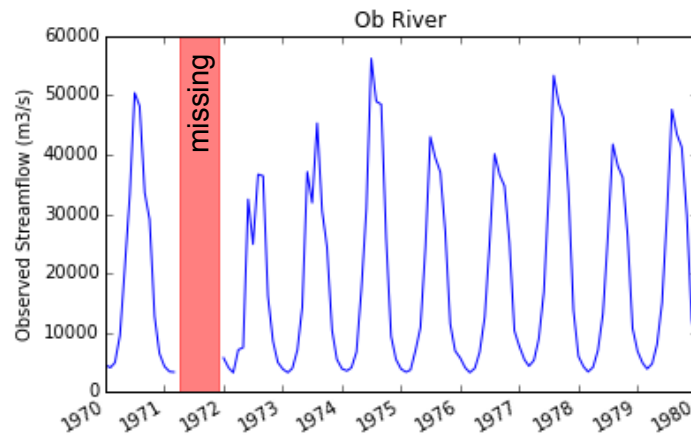

# Geospatial Data

Global FluxNET Observations




You are here

UMD Vegetation Types Per Grid Cell



RASM Surface Air Temperature





# The pandas Data Model



## **Pandas Index:**

The basic object storing axis labels for all pandas objects

## **Pandas Series:**

One-dimensional ndarray with axis labels (including time series).

## **Pandas DataFrame:**

Two-dimensional size-mutable, potentially heterogeneous tabular data structure with labeled axes (rows and columns). Made up of one or more Series.

## **Pandas Panel:**

Represents wide format panel data, stored as 3-dimensional array

# The pandas DataFrame





the  
index

Missing  
data

| station    | Yenisey | Lena         | Ob    | Amur | Mackenzie | Yukon | Pechora     | Nelson | Khatanga | Kolyma |
|------------|---------|--------------|-------|------|-----------|-------|-------------|--------|----------|--------|
| time       |         |              |       |      |           |       |             |        |          |        |
| 1948-08-01 | 9569    | 29675.396484 | NaN   | NaN  | 17100     | NaN   | 2165.750000 | NaN    | NaN      | NaN    |
| 1948-09-01 | 7826    | 36318.761719 | NaN   | NaN  | 9470      | NaN   | 2209.843750 | NaN    | NaN      | NaN    |
| 1948-10-01 | 8435    | 19653.771484 | NaN   | NaN  | 6100      | NaN   | 2253.906250 | NaN    | NaN      | NaN    |
| 1948-11-01 | 10248   | 10117.241211 | NaN   | NaN  | 3440      | NaN   | 2254.562500 | NaN    | NaN      | NaN    |
| 1948-12-01 | 10220   | 4085.000000  | NaN   | NaN  | 1870      | NaN   | 1989.875000 | NaN    | NaN      | NaN    |
| 1949-01-01 | 10628   | 3133.926758  | 3249  | NaN  | 1390      | NaN   | 1999.593750 | NaN    | NaN      | NaN    |
| 1949-02-01 | 12876   | 2539.136963  | 3058  | NaN  | 1820      | NaN   | 1811.015625 | NaN    | NaN      | NaN    |
| 1949-03-01 | 16504   | 1944.476685  | 2516  | NaN  | 1740      | NaN   | 1633.187500 | NaN    | NaN      | NaN    |
| 1949-04-01 | 15826   | 2085.389160  | 2105  | NaN  | 5070      | NaN   | 1548.812500 | NaN    | NaN      | NaN    |
| 1949-05-01 | 12234   | 3486.129883  | 2981  | NaN  | 7270      | NaN   | 1334.093750 | NaN    | NaN      | NaN    |
| 1949-06-01 | 11142   | 5410.482422  | 5883  | NaN  | 19800     | NaN   | 926.875000  | NaN    | NaN      | NaN    |
| 1949-07-01 | 11681   | 13076.666992 | 17810 | NaN  | 29400     | NaN   | 744.625000  | NaN    | NaN      | NaN    |
| 1949-08-01 | 8569    | 28834.199219 | 43997 | NaN  | 13000     | NaN   | 880.875000  | NaN    | NaN      | NaN    |
| 1949-09-01 | 7939    | 34319.695312 | 40283 | NaN  | 6220      | NaN   | 748.906250  | NaN    | NaN      | NaN    |
| 1949-10-01 | 8087    | 23169.228516 | 18745 | NaN  | 11800     | NaN   | 1220.281250 | NaN    | NaN      | NaN    |
| 1949-11-01 | 6949    | 10456.551758 | 9773  | NaN  | NaN       | NaN   | 1697.531250 | NaN    | NaN      | NaN    |

# pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



# xray

## Great

- Fast
- Flexible
- Missing data


## Not so great

- N-Dimensional objects and operations
- Metadata

- Fundamentally N-dimensional
- Leverages pandas for indexing
- Tries to copy pandas API
- netCDF (Data model, IO)
- Metadata
- Out of core computation

# netCDF data model

N-Dimensional Homogeneous (type) arrays



# The xray Dataset

```
In [5]: print(ds)
```

```
<xray.Dataset>
```

```
Dimensions: (nv4: 4, time: 366, x: 275, y: 205, z: 1)
```

```
Coordinates:
```

```
* time      (time) datetime64[ns] 2000-01-01 2000-01-02 2000-01-03 ...
  xc        (y, x) float64 189.2 189.4 189.6 189.7 189.9 190.1 190.2 190.4 ...
  yc        (y, x) float64 16.53 16.78 17.02 17.27 17.51 17.76 18.0 18.25 ...
* z         (z) float32 0.0
* nv4       (nv4) int64 0 1 2 3
* x         (x) int64 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 ...
* y         (y) int64 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 ...
```


```
Data variables:
```

```
xc_bnds     (y, x, nv4) float64 189.3 189.4 189.2 189.0 189.4 189.6 189.3 ...
yc_bnds     (y, x, nv4) float64 16.33 16.58 16.74 16.49 16.58 16.82 16.98 ...
prcp        (time, z, y, x) float64 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 ...
tas         (time, z, y, x) float64 291.6 291.4 291.5 291.5 290.8 290.7 ...
```

```
Attributes:
```

```
CDI: Climate Data Interface version 1.5.4 (http://code.zmaw.de/projects/cdi)
Conventions: CF-1.4
```

# xray's Dataset is a N-dimensional DataFrame



# Serialization and IO

Because nobody likes parsing messy data

| <b>pandas</b>                 | <b>xray</b>                              |
|-------------------------------|--|
| <code>{read, to}_csv()</code> | <code>open_dataset(), to_netcdf()</code> |
| <code>{read, to}_excel</code> | Backends:                                |
| <code>{read, to}_hdf</code>   | netCDF4                                  |
| <code>{read, to}_sql</code>   | scipy                                    |
| <code>{read, to}_json</code>  | h5netcdf                                 |
| <code>{read, to}_html</code>  | pydap                                    |

# Leveraging the index and labels

Selecting, slicing, and subsetting

```
# Xray Style
In [6]: ds.sel(time=slice('2000-01-15', '2000-03-15')).mean(dim='station')

# Numpy Style
In [7]: array[15:60, :, :].mean(axis=2)
```

Label based methods are:

- More reliable
- Easier to maintain
- Easier to read



# Leveraging the index

Resampling: frequency conversion and resampling of regular time-series data.

```
# Resample Dataset from 6-hourly to monthly means
In [8]: ds.resample('MS', dim='time', how='mean')

# Resample Dataset from 6-hourly to daily sums
In [9]: ds.resample('D', dim='time', how='sum')
```

# Leveraging the index

Grouping: split – apply – combine

- **Split** your data into multiple independent groups.
- **Apply** some function to each group.
- **Combine** your groups back into a single data object.

```
# Get Monthly Climatology
```

```
In [10]: ds.groupby('time.month').mean(dim='time')
```

```
# Or Seasonal Climatology
```

```
In [11]: ds.groupby('time.season').mean(dim='time')
```

# Other stuff

```
# Fancy Indexing:
```

```
In [12]: ds.isel_points(lon=[1, 3], lat=[2, 2])
```

```
# Vectorized math and broadcasting (based on dim name)
```

```
In [13]: ds.x * ds.y + ds.z
```


```
# Alignment:
```

```
In [14]: x, y = xray.align(x, y, join='outer')
```

# xray and pandas make working with(out) missing data much easier

```
# Select it:  
In [15]: df.isnull()  
  
# Skip it:  
In [16]: df.mean()  
  
# Drop it:  
In [17]: df.dropna()  
  
# Fill it:  
In [18]: df.fillna()
```

```
# For Example:  
In [19]: ds.Tair.\  
mean(dim='time').plot ()
```



# Interoperability with Numpy/Scipy

```
# Calculate the 5th percentile for each month
```

```
In [20]: percents = ds.groupby('time.month').reduce(np.percentile, dim='time', q=5)
```

```
In [21]: percents
```

```
Out[21]:
```

```
<xray.Dataset>
```

```
Dimensions: (lat: 25, lon: 53, month: 12)
```

```
Coordinates:
```

```
* lat      (lat) float32 75.0 72.5 70.0 67.5 65.0 62.5 60.0 57.5 55.0 52.5 ...
* lon      (lon) float32 200.0 202.5 205.0 207.5 210.0 212.5 215.0 217.5 ...
* month     (month) int64 1 2 3 4 5 6 7 8 9 10 11 12
```

```
Data variables:
```

```
air      (month, lat, lon) float64 237.5 237.8 237.6 237.4 236.7 235.8 ...
```

# Xray and Pandas Interoperability

```
In [22]: df = percents.sel(lat=70., lon=202.5).to_dataframe()
```

```
In [23]: print(df)
```

```
air  lat  lon
```

month

|   |          |    |       |
|---|----------|----|-------|
| 1 | 240.2350 | 70 | 202.5 |
|---|----------|----|-------|

2      239.3000      70      202.5

```
In [24]: xray.Dataset.from_dataframe(df)
```

&lt;xray.Dataset&gt;

Dimensions: (month: 12)

## Coordinates:

```
* month      (month) int64 1 2 3 4 5 6 7 8 9 10 11 12
```

Data variables:

```
air      (month) float64 240.2 239.3 243.6 247.5 257.2 270.5 274.8 272.6 ...
```

```
lat      (month) float32 70.0 70.0 70.0 70.0 70.0 70.0 70.0 70.0 70.0 70.0 ...
```

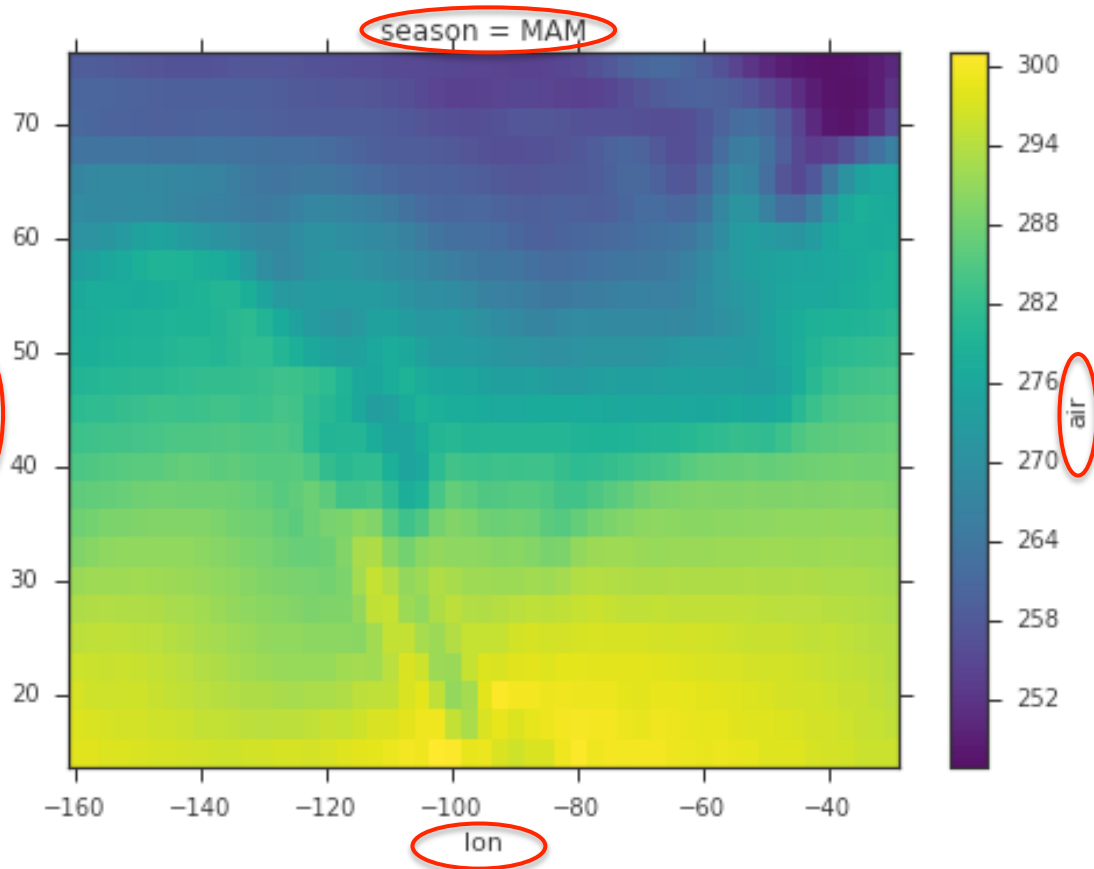
```
lon      (month) float32 202.5 202.5 202.5 202.5 202.5 202.5 202.5 202.5 202.5 ...
```

# Labeled data can plot itself

```
# For Example:  
In [25]: da.plot()
```




People need incentives to  
preserve metadata



# FacetGrid

Borrowed from Seaborn


```
# For Example:  
In [26]:  
da.plot(col='season',  
        , col_wrap=2)
```





# Plotting with Cartopy

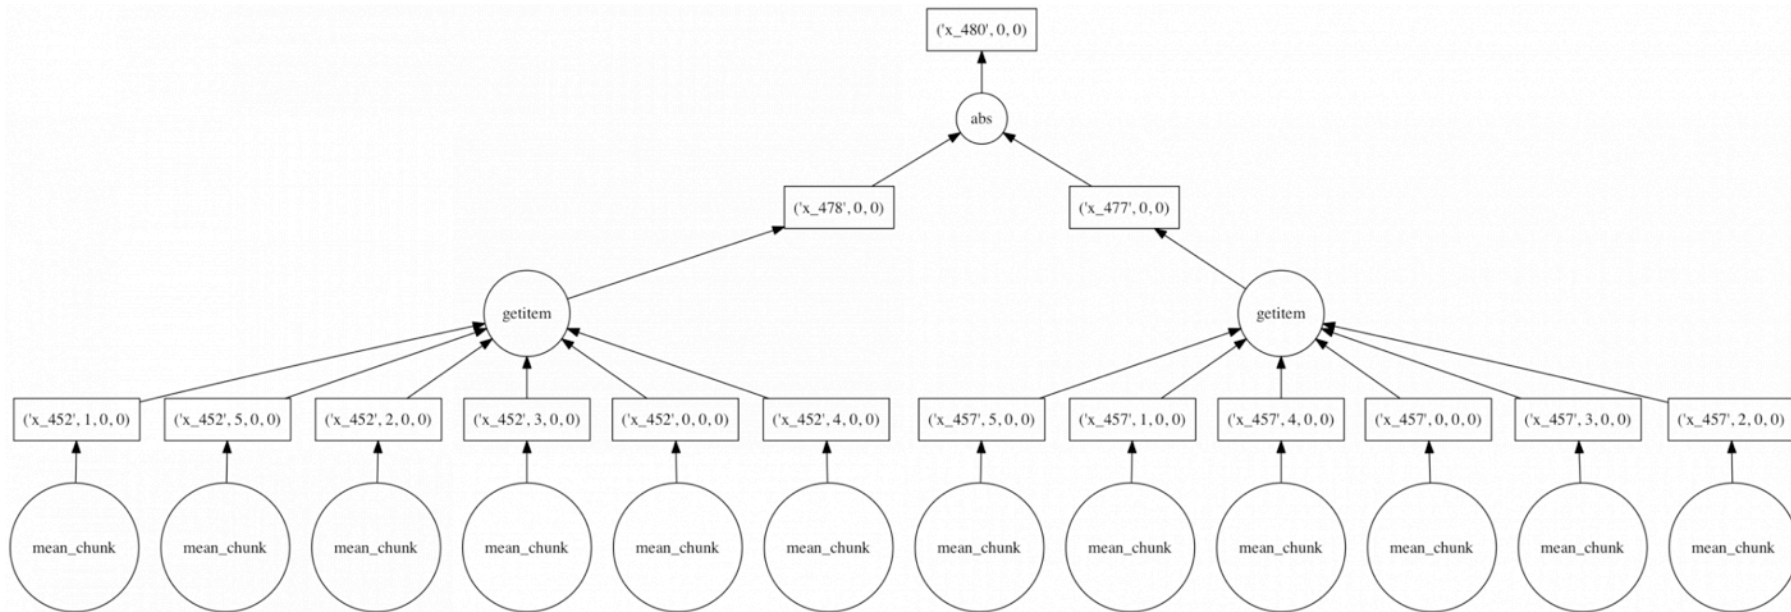

```
ax = plt.axes(projection=ccrs.LambertConformal())  
mappable = da.plot(levels=10, transform=ccrs.PlateCarree())
```




# Getting out of core: Dask

Data processing tasks are often embarrassingly parallel.

Lazy Computation




# Who's in the game



Blaze  
Dask

Dask enables parallel computing through task scheduling and blocked algorithms.

What is medium data?




Cartopy is a Python package designed to make drawing maps for data analysis and visualization as easy as possible.

Replacement? to Matplotlib Basemap.


# Scientific Python's Horizons

We should be able to write new dtypes in Python

Categorical



Missing data



Dates & times



Physical  
Units

$$52.8 \text{ ft/s} \\ = 36 \text{ mi/h}$$

# **We should stop rewriting the same things...**

Working Assumption:

It is better to push repetitive tasks to well tested, optimized packages (e.g. pandas/xray/NumPy)

## **So...**

Don't hide your code

Contribute to open source projects

Build on the success of others

# Questions?

- > Joe Hamman
- > Computational Hydrology, CEE
- > Email: [jhamman1@uw.edu](mailto:jhamman1@uw.edu)
- > Web: [joehamman.com](http://joehamman.com)
- > Github: [jhamman](https://github.com/jhamman)